

Participant Guide & Official Rules

University of Chicago · May 16–17, 2026 · Evaluation through May 31

Welcome

We're thrilled to have you at Prophet Hacks. This is a fast, intense weekend of building AI agents that have to do something genuinely hard: predict the future. Whether you're shipping a forecasting agent or a trading agent, we want you to leave with a working system, new collaborators, and ideally a leaderboard spot.

This document covers everything you need to compete: how check-in works, how your agent will be evaluated, FAQs, and the official rules. Skim the Quick Facts below, then dig into whatever section matters for what you're doing right now.

Quick Facts

Event	Prophet Hacks 2026
Kickoff	Saturday, May 16, 2026 at 9:00 AM CT
Build window	~32 hours (kickoff → Sunday 5:00 PM)
Submission deadline	Sunday, May 17, 2026 at 5:00 PM CT
Evaluation window	2 weeks: May 17 → May 31, 2026
Winners announced	Monday, Jun 1, 2026
Venue	University of Chicago — John Crerar Library
Tracks	Forecasting Track or Trading Track (one per team)
Team size	1–10 people (solo allowed; teaming up encouraged)
Prizes per track	Winning team: a sponsored trip to South Korea (up to \$2,000) to present at the ICML workshop · Runner-up: \$500
Contact	contact@prophetarena.co · Discord

Hackathon Tracks

You can enter one of the following tracks. Pick the one that matches what you most want to build — they reward different things.

Forecasting Track

Build an agent that predicts probabilities for real-world events as accurately as possible. This track rewards careful research, evidence aggregation, and calibration. You don't need a trading strategy to win — you need to be right, consistently.

Trading Track

Build an agent that turns predictions into profitable positions on prediction markets. This track rewards good decision-making under uncertainty: risk management, sizing, and execution often matter as much as raw forecasting skill. You can win here without having the single best forecaster, if your strategy is sharp.

Heads-up: keep an eye on your API costs. During the 2-week evaluation window, your agent will be queried continuously and runs on **your own API keys**. Costs can add up quickly — please be mindful of model choice, token usage, and request volume. If at any point the cost or operational burden becomes too much, you may withdraw your agent from the evaluation; just let the organizers know at contact@prophetarena.co. Withdrawing forfeits your eligibility for that track's prize but is always preferable to running up a bill you didn't sign up for.

Check-In & Getting Started

1. Register your team and pick a track

At check-in, each participant signs up under a team affiliation and the team selects a track. Any team size from 1 to 10 works. If you don't have a team yet, no problem — we'll run a team-formation session right after the Saturday kickoff. Once you've decided on your team and track, fill in the check-in form:

<https://forms.gle/1sz1x1BkYdGuLXmR9>

2. Grab your API credits

After check-in, you may receive a development API key from OpenRouter with a \$50 budget cap to prototype with during the hackathon. Keys are issued first-come, first-served to the first 100 participants and are deactivated when the event ends.

During the evaluation window, teams deploy their agents using **their own API keys** — the dev keys are for prototyping only.

3. Build your agent

Use whatever models, frameworks, and tools you'd like (subject to the Fair Play rules below). The provided starter packages are a fine launchpad, but the core of your submission should be created during the event.

4. Submit before the deadline

Each team submits the following by Sunday, May 17 at 5:00 PM CT:

- **Devpost submission page** — includes all team members and project details: an overview of the agent, architecture and design decisions, key innovations.
 - **Code repository** — a *public* GitHub repository is required. We recommend releasing under MIT or Apache 2.0. Please also include a script with instructions that the organizers can run the agent in a standardized environment. DO NOT commit any proprietary API keys.
 - **Agent deployed for evaluation** — For the Forecasting Track, your agent must expose an OpenAI-compatible API endpoint the evaluation harness can query. For the Trading Track, your agent must integrate with the provided trading harness.
-

How Evaluation Works

Forecasting Track

- Each team submits an HTTP endpoint that accepts a list of forecasting questions and returns a predicted probability for each, within a 10-minute response window.
- Our evaluation server queries your endpoint continuously over a 2-week evaluation window (with up to 200 events in total) and scores each prediction with the realized outcome under the Brier score.
- Please make your agent reliable: we expect a prediction for every forecasting question. Missed questions count against your completion rate.
- **Total score = (team average Brier – market average Brier) × completion rate.** Higher is better, and the formula is set up so a lower (better) Brier than the market produces a positive score.

To win or place as runner-up in this track, a team's total score must be greater than 0 — meaning your agent must, on average, forecast more accurately than the market.

Trading Track

- Each team runs its trading agent against an evaluation slug provided by the ai-prophet SDK.

- Standings are determined by the sum of your rankings across all evaluation metrics (including Sharpe ratio and PnL). Lowest combined rank wins. For example, a team ranked #1 in Sharpe and #4 in PnL has a combined rank of 5, which finishes below a team ranked #2 and #2 (combined rank 4).

To win or place as runner-up in this track, a team's PnL must be positive and must make at least 14 trades in total — meaning your agent should comfortably profit from trading on the market.

Common Questions

What's the difference between the two tracks?

The Forecasting Track is about doing good research and producing well-calibrated forecasts. The Trading Track is about making good decisions on top of forecasts — risk, sizing, timing. A good trading strategy doesn't strictly require the best forecaster, and a great forecaster won't automatically print money without a sensible trading strategy. Different muscles.

What's the team size?

Capped at 10 people. Solo participation is allowed, but teaming up tends to produce stronger agents, and we'll help you find collaborators at the kickoff.

Can I enter both tracks?

No. You'll need to commit to a single track and team affiliation based on your check-in response.

Do I have to use a specific model or API?

No. Use whatever tools you'd like, as long as you stay within the providers' terms of service and the Fair Play rules below. AI coding assistants are explicitly welcome.

Who do I contact for help?

During the event, ping us on Discord or grab an organizer at the venue. For rules questions or anything sensitive, email contact@prophetarena.co or message an organizer on the Discord server.

Official Rules

These rules govern participation in Prophet Hacks. By registering and submitting, you agree to them.

1. Eligibility

- Open to students, researchers, developers, startup teams, and independent builders.
- Participants may compete individually or in teams of up to 10 people. All team members must be registered.
- Participants must be 18 or older at the time of the event. If you're under 18 and would like to attend, contact us in advance at contact@prophetarena.co — we may be able to accommodate you with signed parental/guardian consent and a chaperone.
- Organizers, judges, and anyone with access to evaluation internals or non-public benchmark data are not eligible to win prizes.

2. Teams

- Teams form before or at kickoff. We run a team-formation session right after the Saturday kickoff if you're flying solo.
- Team size is capped at 10. Solo participation is allowed.
- Team rosters lock at the submission deadline.
- A person may not be on more than one team.

3. Dates & Deadlines

- Build window: ~32 hours, Saturday, May 16 (9:00 AM kickoff) through Sunday, May 17 (5:00 PM).
- Submission deadline: Sunday, May 17, 5:00 PM CT. Late submissions are not accepted.
- Evaluation period: agents run for a 2-week evaluation window after the event (May 17 → May 31).
- Winners are announced after leaderboard results are finalized — Monday, June 1, 2026.

4. What to Build

- Each team builds an AI forecasting agent that produces probabilistic predictions about real-world events (technology, economics, geopolitics, science, and similar).
- Agents may combine LLM reasoning, web search and retrieval, evidence aggregation, probabilistic calibration, and market-based or automated trading strategies.
- Any models, frameworks, APIs, and tools may be used, subject to the Fair Play rules below.

- Projects must be built primarily during the hackathon. Public/open-source libraries and the provided starter packages may be used as a starting point, but the core submitted work must be created during the event.

5. Judging & Evaluation

- **Forecasting Track** — evaluated on forecasting accuracy on the Prophet Arena benchmark.
- **Trading Track** — evaluated on live paper-trading profit on prediction markets during the evaluation window, using the provided harness.
- Teams may only enter one track. Evaluation runs on prophethacks.com with a live leaderboard. Final standings are determined after the 2-week evaluation period.
- Organizers' and judges' decisions are final.
- **Forecasting Track minimum.** To win or place as runner-up, a team's total score must be greater than 0 — meaning the agent must, on average, forecast more accurately than the market.
- **Trading Track minimum.** To win or place as runner-up, a team's PnL must be positive and the agent must execute at least 14 trades during the evaluation window.
- **Right to withdraw.** A team may withdraw its agent from the evaluation at any time by emailing contact@prophetarena.co. Withdrawing forfeits eligibility for that track's prize but releases the team from any further evaluation obligation.
- **Prizes** — for each track, the winning team receives a sponsored trip to South Korea (value up to \$2,000) to present at the ICML workshop, and the runner-up receives \$500. Prizes are awarded per team; teams are responsible for deciding who to attend the trip or how to split any cash award among members.

6. Intellectual Property & Ownership

- Teams retain ownership of the IP in the work they create during the hackathon.
- By submitting, teams grant the organizers a non-exclusive, royalty-free license to review, run, evaluate, and showcase the submission (including on the leaderboard and in event materials) for purposes related to the event.
- Teams are responsible for ensuring they have the rights to all components, data, and dependencies used in their submission.
- A public GitHub repository is required. We recommend an open-source license such as MIT or Apache 2.0, but teams may choose any OSI-approved license that permits the organizers to evaluate and showcase the submission.
- Any pull request submitted to ai-prophet / ai-prophet-core is governed by that project's existing license and contribution terms.

- Use of submissions beyond evaluation and showcasing (e.g., incorporation into Prophet Arena) requires the team's separate consent.

7. Fair Play & Evaluation Integrity

Participants must not:

- Submit previously built projects as new work.
- Manipulate, game, or attempt to reverse-engineer the leaderboard or evaluation harness.
- Access non-public benchmark questions, answers, or resolution data, or use information not available to all participants.
- Train, fine-tune, or condition agents on the specific evaluation questions or their outcomes.
- Violate the terms of service of any API, model, dataset, or tool provider.
- Interfere with other teams' work, infrastructure, or submissions.
- Misrepresent authorship, team membership, or results.

AI tools (coding assistants, LLM agents, etc.) are allowed and encouraged for building — the restriction is on gaming the evaluation, not on using AI to build. Violations may result in disqualification at the organizers' discretion.

A note on changes during the evaluation period: teams may make minimal changes to their agents during the 2-week evaluation window to fix bugs or address reliability issues. This is not encouraged — changes mid-evaluation can introduce inconsistency and may worsen your evaluation results. Keep edits surgical, and document anything you change.

8. Code of Conduct

- Prophet Hacks is committed to a harassment-free, inclusive experience for everyone — in person and on Discord — regardless of background or identity.
- **Expected behavior:** be respectful, collaborative, and professional; assume good faith; keep shared spaces and channels welcoming.
- **Unacceptable behavior:** harassment, discrimination, intimidation, unwelcome attention, hate speech, or disruptive conduct — in talks, workspaces, Discord, or social activities.
- **Reporting:** please contact the organizers immediately at contact@prophetarena.co, or message an organizer directly on Discord or on-site. Reports are handled as discreetly as possible.
- **Enforcement:** organizers may take any action they deem appropriate, including warnings, removal from the event or Discord, and disqualification, without refund of any travel or other costs.
- Applies to all participants, judges, mentors, sponsors, speakers, and organizers.

9. General

- Organizers may update these rules before or during the event; material changes will be announced on Discord.
- Organizers may disqualify any participant or team for rule violations.
- Participation is at your own risk; follow venue (University of Chicago / John Crerar Library) policies and applicable laws.
- Prizes are subject to applicable taxes and any sponsor- or organizer-specified conditions.
- Questions about these rules: **contact@prophetarena.co**.

Build something audacious. We can't wait to see what your agents predict.

— *The Prophet Hacks Organizers*